CS 639: Deep Learning for Computer Vision, Spring 2024
Problem Set 3

Instructor: Yong Jae Lee (yongjaelee@cs.wisc.edu)
TA: Zhuoran Yu (zhuoran.yu@wisc.edu)
TA: Zeyi Huang (zhuang479@wisc.edu)

Due: Wednesday, April 10th, 11:59 PM

Instructions

1. Download the zipped assignment file from Canvas.

2. Once you unzip the downloaded content, please upload the folder to your Google Drive. Then, open each `*.ipynb` notebook file with Google Colab by right-clicking the `*.ipynb` file. No installation or setup is required. For more information, please see this tutorial on using Colab.

3. Next, we recommend editing your `*.py` file on Google Colab, set the ipython notebook and the code side by side. Work through the notebook, executing cells and implementing the codes in the `*.py` file as indicated. You can save your work, both `*.ipynb` and `*.py`, in Google Drive (click "File" -> "Save" or press "Ctrl/Cmd + s") and resume later if you don't want to complete it all at once.

4. While working on the assignment, keep the following in mind:

   - The notebook and the python file have clearly marked blocks where you are expected to write code. **Do not write or modify any code outside of these blocks**.
   - **Do not add or delete cells from the notebook**. You may add new cells to perform scratch computations, but you should delete them before submitting your work.
   - **Run all cells, and do not clear out the outputs, before submitting**. You will only get credit for code that has been run.

5. Once you have completed a notebook, download the completed `uniqueid_PS3.zip` file, which is generated from your last cell of the `fcos.ipynb` file. Submit this to Canvas. Note that only one person from the group will need to do this.

6. You may complete the assignment individually or with a partner (i.e., maximum group of 2 people). If you worked with a partner, provide the name of your partner in the `*.ipynb` file. **We will be using MOSS to check instances of plagiarism/cheating.**

The goal of this assignment is for you to learn how to build an object detection neural network so that you can
- Develop a deep understanding of the object detection pipeline.
- Gain experience using the feature pyramid network for the center-based single stage object detector.

The notebook `fcos.ipynb` will walk you through those pipelines in PyTorch. You are required to write code on `fcos.py`.[1]

[OPTIONAL] Extra credit short answer problems [up to 0.5% each added to final class score]

Save your answers into a PDF, and submit it together with your code to Canvas.

1. Explain why Fast-RCNN is faster than R-CNN, and then why Faster-RCNN is faster than Fast-RCNN.

2. Explain two reasons why we might want to use a fully convolutional network for semantic segmentation.

---

[1] The coding assignment is adapted from Stanford CS 231n.